

 [main](#) ▾



[Phase\\_3\\_Project](#) / [Project 3 Final Notebook.ipynb](#)



**petercvuong** Ran final notebook again to minimize output of some cells

 **History**

 1 contributor

1.12 MB





Centers for Disease  
Control and Prevention

## Contents

- 1. Business Understanding
- 1. EDA and Data Cleaning
- 1. Modeling
  - 3.1 Dummy Model
  - 3.2 DecisionTree Classifier
  - 3.3 LogisticRegression Classifier
- 1. Results
- 1. Conclusion
- 1. Recommendations
- 1. Future Considerations

## 1. Business Understanding

### Tackling Flu and Vaccine Misinformation for the CDC

The mission statement of the CDC is **"to promote health and quality of life by preventing and controlling disease, injury, and disability."**

And so, an integral part of fulfilling this mission is to provide the population with the appropriate information that is needed in order to make the best informed decisions for their health.

Disease and vaccine misinformation are a major

hurdle for the CDC. Lack of information or knowledge of the wrong information can lead to misinformed decisions that will compromise not only a single individual's health and safety but the population's as well.

The main goal of our project is to predict what features are most influential in determining a respondent's knowledge of the H1N1 flu and vaccine. Once we identify what features are most influential to a respondent's knowledge, we can focus on these features as areas of that the CDC should focus on in order to minimize the dangers of misinformation.

## 2. EDA and Data Cleaning

### Data and Limitations

The data that we will be working with in this project comes from the [DrivenData](#) website. This data is a survey response from 2009.

Since this data is from a survey response, it does not give an accurate reflection of the situation and community at the time. The data also has a pretty obvious class imbalance and some clear biases.

### Loading in the Data and some Initial EDA

We are looking at the vaccine data from 2009 about H1N1 flu and vaccine awareness.

The data survey response data that we got from DrivenData had about 27,000 recorded responses.

```
In [1]: # Importing the appropriate libraries that will be used
import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
from sklearn.feature_selection import RFE
from sklearn.preprocessing import PolynomialFeatures
import statsmodels
```

```

from statsmodels.formula.api import ols
from sklearn.dummy import DummyRegressor

from scipy import stats
from sklearn.preprocessing import OneHotEncoder
from folium.plugins import FastMarkerCluster

from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegressor
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, recall_score
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import roc_auc_score, plot_roc_curve

```

In [2]:

```

# Loading in the pre-split datasets that were given
vaccinetrainingdf = pd.read_csv("data/training_set_features.csv")
vaccinetestdf = pd.read_csv("data/test_set_features.csv")
vaccinelabelsdf = pd.read_csv("data/training_set_labels.csv")

```

In [3]:

```

# Initial checking to see what data types we are working with
# Commenting outputs out from notebook to reduce size
## vaccinetrainingdf.info()

```

As we can see there are some missing values, and some of the values are objects. Since some of the values in these columns are objects, we know that we have to one hot encode the values in order to implement them mathematically into our models.

In [4]:

```

vaccinetrainingdf.head()

```

Out[4]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavior
0	0	1.0	0.0	
1	1	3.0	2.0	
2	2	1.0	1.0	
3	3	1.0	1.0	
4	4	2.0	1.0	

5 rows × 36 columns

```
In [5]: # We see here that the data is already split almost
# Commenting outputs out from notebook to reduce
## vaccinetestdf.info()
```

```
In [6]: # Utilizing this block of code just to display all
pd.set_option('max_columns', None)
vaccinetestdf.head()
```

```
Out[6]:
```

	respondent_id	h1n1_concern	h1n1_knowledge	behavior
--	---------------	--------------	----------------	----------

0	26707	2.0	2.0	
---	-------	-----	-----	--

1	26708	1.0	1.0	
---	-------	-----	-----	--

2	26709	2.0	2.0	
---	-------	-----	-----	--

3	26710	1.0	1.0	
---	-------	-----	-----	--

4	26711	3.0	1.0	
---	-------	-----	-----	--

```
In [7]: vaccinelabelsdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   respondent_id    26707 non-null  int64
1   h1n1_vaccine     26707 non-null  int64
2   seasonal_vaccine 26707 non-null  int64
dtypes: int64(3)
memory usage: 626.1 KB
```

Next, we're going to be doing some data cleaning to get rid of any extraneous columns that we determined were not relevant to our business problem of tackling misinformation in flu and vaccine awareness.

```
In [8]: # Dropping data we deemed unnecessary and irrelevant
columns_to_drop = ['respondent_id', 'h1n1_knowledge',
                    'employment_status', 'rent_or_ownership',
                    'health_worker']
X_train = vaccinetrainingdf.copy().drop(columns_to_drop)
X_test = vaccinetestdf.copy().drop(columns_to_drop)
# Setting the training and test data just for the
```

```
# Setting the y_train and y_test to just be the  
y_train = vaccinetraindf['h1n1_knowledge']  
y_test = vaccinetestdf['h1n1_knowledge']
```

In [9]:

```
# Checking our data again to see that we dropped  
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 26707 entries, 0 to 26706  
Data columns (total 20 columns):  
#   Column                                Non-Null Count  
Dtype  
---  ---  
0    h1n1_concern                        26615 non-null  
float64  
1    behavioral_antiviral_meds          26636 non-null  
float64  
2    behavioral_avoidance               26499 non-null  
float64  
3    behavioral_face_mask               26688 non-null  
float64  
4    behavioral_wash_hands              26665 non-null  
float64  
5    behavioral_large_gatherings        26620 non-null  
float64  
6    behavioral_outside_home            26625 non-null  
float64  
7    behavioral_touch_face              26579 non-null  
float64  
8    doctor_recc_h1n1                  24547 non-null  
float64  
9    doctor_recc_seasonal               24547 non-null  
float64  
10   chronic_med_condition              25736 non-null  
float64  
11   opinion_h1n1_vacc_effective          26316 non-null  
float64  
12   opinion_h1n1_risk                    26319 non-null  
float64  
13   opinion_h1n1_sick_from_vacc          26312 non-null  
float64  
14   opinion_seas_vacc_effective          26245 non-null  
float64  
15   age_group                           26707 non-null  
object  
16   education                           25300 non-null  
object  
17   race                                26707 non-null  
object  
18   sex                                 26707 non-null  
object  
19   income_poverty                      22284 non-null  
object  
dtypes: float64(15), object(5)  
memory usage: 4.1+ MB
```

In [10]:

```
X_test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26708 entries, 0 to 26707
Data columns (total 20 columns):
 #   Column                                Non-Null Count
Dtype
---  -
0   h1n1_concern                        26623 non-null
float64
1   behavioral_antiviral_meds          26629 non-null
float64
2   behavioral_avoidance               26495 non-null
float64
3   behavioral_face_mask               26689 non-null
float64
4   behavioral_wash_hands              26668 non-null
float64
5   behavioral_large_gatherings        26636 non-null
float64
6   behavioral_outside_home            26626 non-null
float64
7   behavioral_touch_face              26580 non-null
float64
8   doctor_recc_h1n1                  24548 non-null
float64
9   doctor_recc_seasonal               24548 non-null
float64
10  chronic_med_condition              25776 non-null
float64
11  opinion_h1n1_vacc_effective         26310 non-null
float64
12  opinion_h1n1_risk                   26328 non-null
float64
13  opinion_h1n1_sick_from_vacc         26333 non-null
float64
14  opinion_seas_vacc_effective         26256 non-null
float64
15  age_group                          26708 non-null
object
16  education                          25301 non-null
object
17  race                               26708 non-null
object
18  sex                                26708 non-null
object
19  income_poverty                     22211 non-null
object
dtypes: float64(15), object(5)
memory usage: 4.1+ MB

```

## Feature Engineering

Created some frequently used functions that we will be utilizing throughout our project

```

In [11]: # Defined a OneHotEncoder function for ease of use
def OHE(X_train, categories):

```

```

onehot = OneHotEncoder(sparse=False, handle_unknown='ignore')
x_train_cat = pd.DataFrame(onehot.fit_transform(x_train_cat).toarray())
x_train_cat.columns = onehot.get_feature_names_out()

# Reset indices to avoid merging conflicts
x_train_cat.reset_index(drop=True, inplace=True)
X_train.reset_index(drop=True, inplace=True)

# Joined the OHE dataframe to the dataframe
x_train_df = X_train.drop(columns=['category'], axis=1).join(x_train_cat)
return x_train_df

# Defined a function that takes in parameters to
def confusion_and_metrics(model, X_test, y_test):
    # Accuracy Score
    print(f"Accuracy Score: {model.score(X_test, y_test)}")

    # Precision Score
    print(f"Precision Score: {precision_score(y_test, model.predict(X_test))}")

    # Plot confusion matrix for visualization
    plot_confusion_matrix(model, X_test, y_test)

# Defined a function to take in column name and dataframe
def print_odds(dataframe, column_name):
    # Prints out the name of the column and it's data type
    print(f"{column_name}: {dataframe[column_name].dtypes}")

    # Prints out the odds value of the column
    print(f"Odds: {np.exp(dataframe[column_name].mean())}")

```

In [12]:

```
# X_train
```

## SimpleImputer to Account for NaN Values

Prior to running some classification models on our data, we looked at it again and noticed that there were still a couple missing values.

In order to rectify this, we created a simple imputer to replace the NaN values with the most frequent value (otherwise known as the mode) in its respective column.

We chose to use the mode to replace these NaN values because using the mode will keep the distribution of the data consistent.

In [13]:

```

# Created a SimpleImputer to replace the NaN values
from sklearn.impute import SimpleImputer
imputer = SimpleImputer(strategy='most_frequent')
imputed_X_train = imputer.fit_transform(X_train)

```



```
imputed_X_train_df = pd.DataFrame(imputed_X_train)
```

```
In [14]: imputed_X_test = imputer.transform(X_test)
imputed_X_test_df = pd.DataFrame(imputed_X_test)
imputed_X_test_df
```

```
Out[14]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13
0	2	0	1	0	1	1	0	1	0	0	0	5	1	1
1	1	0	0	0	0	0	0	0	0	0	0	4	1	1
2	2	0	0	1	1	1	1	1	0	0	0	5	4	2
3	1	0	0	0	0	0	0	0	1	1	1	4	2	2
4	3	1	1	0	1	1	1	1	0	0	0	5	2	4
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
26703	1	0	1	0	1	0	0	1	1	1	0	4	2	2
26704	3	0	1	0	1	1	1	1	0	0	0	4	1	1
26705	0	0	0	0	0	0	0	0	0	0	0	4	3	1
26706	3	0	1	0	1	0	1	0	0	0	0	2	3	4
26707	2	0	0	0	1	0	0	1	1	0	0	5	1	2

26708 rows × 20 columns



```
In [15]: # After doing the imputation and renaming, check
imputed_X_test_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26708 entries, 0 to 26707
Data columns (total 20 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0    0          26708 non-null  object
 1    1          26708 non-null  object
 2    2          26708 non-null  object
 3    3          26708 non-null  object
 4    4          26708 non-null  object
 5    5          26708 non-null  object
 6    6          26708 non-null  object
 7    7          26708 non-null  object
 8    8          26708 non-null  object
 9    9          26708 non-null  object
10   10         26708 non-null  object
11   11         26708 non-null  object
12   12         26708 non-null  object
13   13         26708 non-null  object
14   14         26708 non-null  object
15   15         26708 non-null  object
16   16         26708 non-null  object
17   17         26708 non-null  object
18   18         26708 non-null  object
19   19         26708 non-null  object
dtypes: object(20)
memory usage: 4.1+ MB

```

After imputing the data, we recognize that the column names have disappeared but the indices are still there. To resolve this, we create a dictionary with the original column names and call the rename function on this new data frame's columns.

```

In [16]: # Extracting column names into a dictionary
dictionary_of_names = {columns: index for index,

# Flipping the column keys and values
dictionary_of_names_flipped = {dictionary_of_name

# Checking to see if the column names were extra
dictionary_of_names_flipped

```

```

Out[16]: {0: 'h1n1_concern',
1: 'behavioral_antiviral_meds',
2: 'behavioral_avoidance',
3: 'behavioral_face_mask',
4: 'behavioral_wash_hands',
5: 'behavioral_large_gatherings',
6: 'behavioral_outside_home',
7: 'behavioral_touch_face',
8: 'doctor_recc_h1n1',
9: 'doctor_recc_seasonal',
10: 'chronic_med_condition',
11: 'opinion_h1n1_vacc_effective',
12: 'opinion_h1n1_risk',
13: 'opinion_h1n1_risk',
14: 'opinion_h1n1_risk',
15: 'opinion_h1n1_risk',
16: 'opinion_h1n1_risk',
17: 'opinion_h1n1_risk',
18: 'opinion_h1n1_risk',
19: 'opinion_h1n1_risk',
20: 'opinion_h1n1_risk'}

```

```
13: 'opinion_h1n1_sick_from_vacc',
14: 'opinion_seas_vacc_effective',
15: 'age_group',
16: 'education',
17: 'race',
18: 'sex',
19: 'income_poverty'}
```

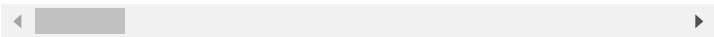
In [17]: *# Created new variables for the training and test*  
imputed\_X\_train\_df\_plus\_column\_names = imputed\_X\_  
imputed\_X\_test\_df\_plus\_column\_names = imputed\_X\_

In [18]: *# Calling the new dataframe variables to check if*  
imputed\_X\_test\_df\_plus\_column\_names

Out[18]:

	h1n1_concern	behavioral_antiviral_meds	behaviora
0	2		0
1	1		0
2	2		0
3	1		0
4	3		1
...	...		...
26703	1		0
26704	3		0
26705	0		0
26706	3		0
26707	2		0

26708 rows × 20 columns



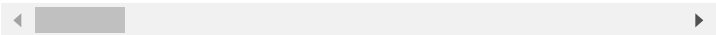
In [19]: imputed\_X\_train\_df\_plus\_column\_names

Out[19]:

	h1n1_concern	behavioral_antiviral_meds	behaviora
0	1		0

1	3	0
2	1	0
3	1	0
4	2	0
...	...	...
26702	2	0
26703	1	0
26704	2	0
26705	1	0
26706	0	0

26707 rows × 20 columns

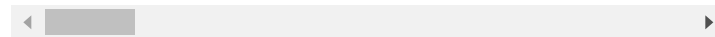


```
In [20]: X_train
```

Out[20]:

	h1n1_concern	behavioral_antiviral_meds	behaviora
0	1.0	0.0	
1	3.0	0.0	
2	1.0	0.0	
3	1.0	0.0	
4	2.0	0.0	
...	...	...	
26702	2.0	0.0	
26703	1.0	0.0	
26704	2.0	0.0	
26705	1.0	0.0	
26706	0.0	0.0	

26707 rows × 20 columns



After doing some initial data cleaning and making sure that our data was uniform, we next want to address the problem of having the `object` type in our columns. In order to address this, we apply a `OneHotEncoder` to these object columns of `age_group`, `education`, `race`, `sex`, and `income_poverty`.

We check the values in each of these object columns to see how many variables will be OneHotEncoded.

```
In [21]: X_train['age_group'].value_counts()
```

```
Out[21]: 65+ Years      6843
55 - 64 Years    5563
45 - 54 Years    5238
18 - 34 Years    5215
35 - 44 Years    3848
Name: age_group, dtype: int64
```

```
In [22]: X_train['education'].value_counts()
```

```
Out[22]: College Graduate    10097
Some College      7043
12 Years          5797
< 12 Years        2363
Name: education, dtype: int64
```

```
In [23]: X_train['race'].value_counts()
```

```
Out[23]: White      21222
Black      2118
Hispanic   1755
Other or Multiple  1612
Name: race, dtype: int64
```

```
In [24]: X_train['sex'].value_counts()
```

```
Out[24]: Female    15858
Male      10849
Name: sex, dtype: int64
```

```
In [25]: X_train['income_poverty'].value_counts()
```

```
Out[25]: <= $75,000, Above Poverty    12777
> $75,000                        6810
Below Poverty                    2697
Name: income_poverty, dtype: int64
```

```
name: income_poverty, dtype: int64
```

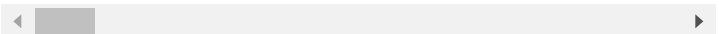
After counting the values, we see that **18** columns will be added

```
In [26]: # Called the OHE function we made and assigned new
ohe_training_df = OHE(imputed_X_train_df_plus_columns)
ohe_test_df = OHE(imputed_X_test_df_plus_columns)
ohe_training_df
```

```
Out[26]:
```

	h1n1_concern	behavioral_antiviral_meds	behavioral_risk_factors
0	1	0	0
1	3	0	0
2	1	0	0
3	1	0	0
4	2	0	0
...	...	...	...
26702	2	0	0
26703	1	0	0
26704	2	0	0
26705	1	0	0
26706	0	0	0

26707 rows × 33 columns



We finished OneHotEncoding the object values and now we have to bin the target values.

Based on the [data dictionary](#), our target values reside in the `h1n1_knowledge` column where the responses are recorded as such:

- 0 = No knowledge
- 1 = A Little Knowledge
- 2 = A Lot of Knowledge

For our project, we are going to bin the 0s and 1s together because those who respond as having little to no knowledge of the flu and vaccine are most prone to misinformation.

We will then be turning all the 2 responses into 1s so that we have a simple binary categorization where:

- 0 = Little/No Knowledge
- 1 = A Lot of Knowledge

```
In [27]: # Instead of calling SimpleImputer and removing NaN values  
# imputation which replaced all the NaN values with 0.0  
# which in this case would be 1.0 (Little knowledge)  
y_train.replace(np.nan, 1.0, inplace = True)  
y_test.replace(np.nan, 1.0, inplace = True)
```

```
In [28]: # Checking to see if we replaced the NaN values correctly  
y_test.isna().value_counts()
```

```
Out[28]: False      26708  
Name: h1n1_knowledge, dtype: int64
```

```
In [29]: y_train.isna().value_counts()
```

```
Out[29]: False      26707  
Name: h1n1_knowledge, dtype: int64
```

```
In [30]: # Binning all the 1.0s with the 0.0s  
y_train.replace(1.0, 0.0, inplace = True)  
y_test.replace(1.0, 0.0, inplace = True)  
# Replacing all the 2.0s with 1.0s  
y_train.replace(2.0, 1.0, inplace = True)  
y_test.replace(2.0, 1.0, inplace = True)
```

```
In [31]: # Checking to see if we replaced our values correctly  
y_train.value_counts()
```

```
Out[31]: 0.0      17220  
        1.0       9487  
Name: h1n1_knowledge, dtype: int64
```

```
In [32]: y_test.value_counts()
```

```
Out[32]: 0.0      17193  
        1.0       9515  
Name: h1n1_knowledge, dtype: int64
```

```
In [33]: # Commenting outputs out from notebook to reduce  
# ohe_training_df.info()
```

## SMOTE for Class Imbalance

After we bin our target and features together, we recognize that our target class is severely imbalanced.

To address this class imbalance, we implement SMOTE to undersample our 0 class.

```
In [34]: # Since our data is severely imbalanced, we utilize SMOTE
# Since we SMOTE our training dataset, we must SMOTE the test set as well

from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
from collections import Counter

over = SMOTE(sampling_strategy=0.7)
under = RandomUnderSampler(sampling_strategy=0.8)

X_smote, y_smote = over.fit_resample(X_train, y_train)
X_test_smote, y_test_smote = under.fit_resample(X_test, y_test)

counter = Counter(y_train)
test_counter = Counter(y_test_smote)
print(counter)
print(test_counter)
```

```
Counter({0.0: 17220, 1.0: 9487})
Counter({0.0: 11893, 1.0: 9515})
```

## Checking for Preliminary Feature Importance

For the final part of our EDA and Data cleaning, we want to check and see what features are seemingly most important to our respondents.

```
In [35]: # Feature columns
X = ohe_training_df.iloc[:,0:33]
# Target column - H1N1 Knowledge
y = vaccinetraingdf.iloc[:,2]

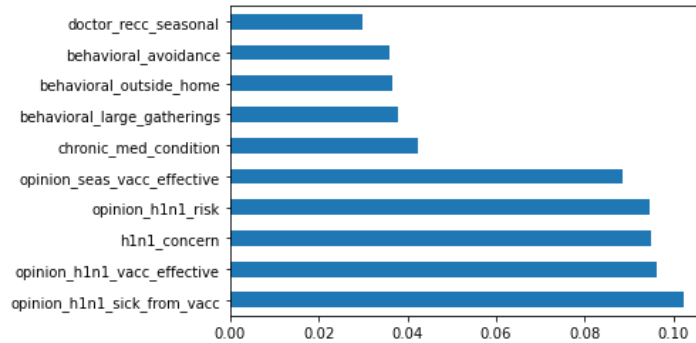
from sklearn.ensemble import ExtraTreesClassifier
import matplotlib.pyplot as plt

# Instantiate model
modelfeatures = ExtraTreesClassifier()
modelfeatures.fit(X,y)
print(modelfeatures.feature_importances_) # use l
# Plot graph of feature importances for better v
feat_importances = pd.Series(modelfeatures.feature
feat_importances.nlargest(10).plot(kind='barh')
plt.show()
```

```
[0.09506911 0.01447003 0.03583344 0.01769623 0.02
254268 0.03772919
 0.03655344 0.02747728 0.02085327 0.02988187 0.04
221946 0.09611108
 0.09469669 0.10224987 0.08850742 0.01540307 0.01
607095 0.01713853
 0.01701695 0.01558146 0.01093478 0.01123703 0.02
000000 0.00000000]
```



```
630258 0.00947193
0.00769741 0.0072257 0.00790023 0.01453193 0.01
239084 0.01235294
0.01141553 0.01665281 0.0087843 ]
```



We notice that the top 5 features that are most important to survey respondents are:

- opinion\_h1n1\_sick\_from\_vacc
- opinion\_h1n1\_vacc\_effective
- h1n1\_concern
- opinion\_h1n1\_risk
- opinion\_seas\_vacc\_effective

## 3. Modeling

Now that our data is cleaned, we can go into the modeling.

### 3.1 Model 1 (Dummy Classifier)

First we want to create a DummyClassifier model that will serve as the baseline for our model performance comparison. A DummyClassifier model in this case would mean that based on the given data, the dummy model would correctly identify our predictions **50%** of the time.

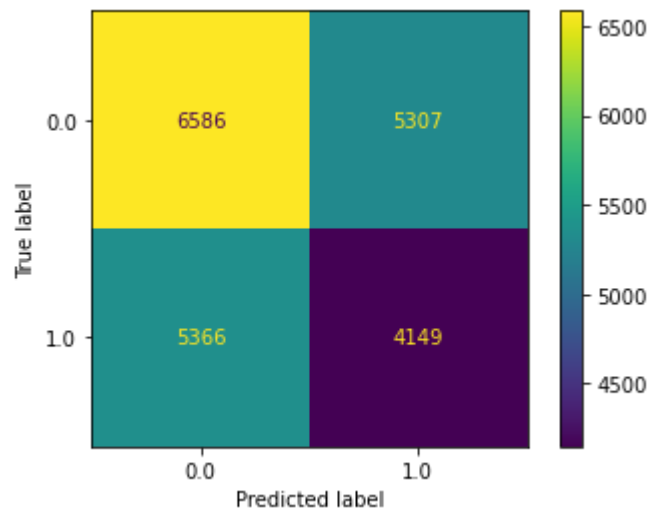
```
In [36]: # Created Dummy Classifier model to look at simple
from sklearn.dummy import DummyClassifier
dummy = DummyClassifier()
dummy.fit(X_smote, y_smote)
y_pred = dummy.predict(X_smote)
y_test_pred = dummy.predict(X_test_smote)
y_pred_df = pd.DataFrame(y_pred)
dummy.score(X_test_smote, y_test_smote)
```

C:\Users\Beter\anaconda3\envs\learn-env\lib\site-packages\sklearn\dummy.py:131: FutureWarning: The

```
packages (sklearn.dummy.py:131: FutureWarning: The
default value of strategy will change from strati
fied to prior in 0.24.
  warnings.warn("The default value of strategy wi
ll change from "
0.5091554559043349
Out[36]:
```

```
In [37]: # Called function and printed out confusion and
confusion_and_metrics(dummy, X_test_smote, y_test)
```

Accuracy Score: 0.508  
Precision Score: 0.442



We print out the `dummy.score` to see that the accuracy score is about 50%, just as we expected.

We call our `confusion_and_metrics` function that we defined above in order to produce the evaluation metrics of Accuracy and Precision and a confusion matrix for easier visualization.

## 3.2 Model 2 (Decision Tree Classifier)

Next we will create an inferential DecisionTree Classifier in order to identify our **most important features**. After we identify our most important features, we can then run a classifying LogisticRegression model to measure our predictions on the dataset.

```
In [38]: # Displayed cross validation score for the dummy
from sklearn.model_selection import train test si
```

```
dummy_cross_val = cross_val_score(dummy, X_smote, y_smote, cv=5)
dummy_cross_val
```

C:\Users\Beter\anaconda3\envs\learn-env\lib\site-packages\sklearn\dummy.py:131: FutureWarning: The default value of strategy will change from stratified to prior in 0.24.

```
warnings.warn("The default value of strategy will change from "
Out[38]: array([0.50480206, 0.51628016, 0.50995549, 0.50503631, 0.51253221])
```

```
In [39]: from sklearn.tree import DecisionTreeClassifier,
from sklearn.model_selection import GridSearchCV

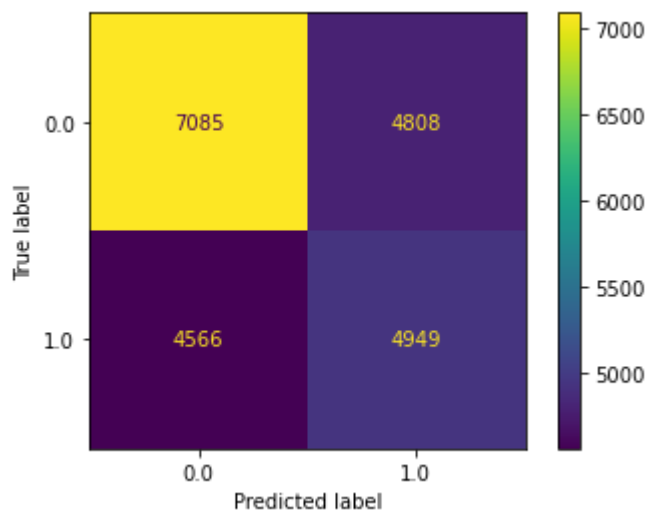
dt = DecisionTreeClassifier (random_state = 10)
dt.fit(X_smote, y_smote)
y_dt_pred = dt.predict(X_smote)
y_dt_test_pred = dt.predict(X_test_smote)
dt.score(X_test_smote, y_test_smote)
```

Out[39]: 0.562126307922272

Similarly like how we checked the DummyClassifier, we also check the plot the confusion matrix and check the metrics of our baseline decision tree.

```
In [40]: # Called function and printed out confusion matrix
confusion_and_metrics(dt, X_test_smote, y_test_smote)
```

Accuracy Score: 0.562  
Precision Score: 0.507



After running our initial DecisionTree Classifier, we got an accuracy score of about **56%**. As you can see, this accuracy score is only about 6% better than the baseline.

## Implementing GridSearchCV to Find Optimal Hyperparameters

We decided to use a `GridSearchCV` in order find the best hyperparameters to pass into our `DecisionTree Classifier` so that we can find the most important features to focus on.

```
In [41]: # Created grid parameter to perform a GridSearch
grid = {
    'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'min_samples_split': [100, 500, 1000, 5000],
    'min_samples_leaf': [100, 500, 1000, 5000]
}
# initializing our grid search with the grid parameter
gs = GridSearchCV(estimator = dt, param_grid = grid)
gs.fit(ohe_training_df, y_train)
```

```
Out[41]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(
    random_state=10),
    param_grid={'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'min_samples_leaf': [100, 500, 1000, 5000],
    'min_samples_split': [100, 500, 1000, 5000]})
```

After we run our `GridSearchCV`, we print out the `best_params_`, `best_score_`, and the `best_estimator_` to get the optimal parameters and metrics based on the grid search results.

```
In [42]: gs.best_params_
```

```
Out[42]: {'max_depth': 7, 'min_samples_leaf': 100, 'min_samples_split': 100}
```

```
In [43]: gs.best_score_
```

```
Out[43]: 0.6770884108866996
```

```
In [44]: gs.best_estimator_.score(ohe_test_df, y_test)
```

```
Out[44]: 0.6696495432080276
```

```
In [45]: gs.n_features_in_
```

```
Out[45]: 33
```

```
In [46]: # Converted the results from the GridSearch to a
```

```
# converted the results from the GridSearch to a
pd.DataFrame(gs.cv_results_)
```

Out[46]:

	mean_fit_time	std_fit_time	mean_score_time	std_sco
--	---------------	--------------	-----------------	---------

0	0.052703	0.004094	0.010736	C
---	----------	----------	----------	---

1	0.047641	0.001899	0.010346	C
---	----------	----------	----------	---

2	0.051506	0.002567	0.011032	C
---	----------	----------	----------	---

3	0.055478	0.012545	0.010883	C
---	----------	----------	----------	---

4	0.051519	0.003832	0.011048	C
---	----------	----------	----------	---

...	...	...	...	
-----	-----	-----	-----	--

155	0.055637	0.002057	0.009565	C
-----	----------	----------	----------	---

156	0.050849	0.002426	0.009955	C
-----	----------	----------	----------	---

157	0.049724	0.003120	0.009994	C
-----	----------	----------	----------	---

158	0.047685	0.001417	0.009380	C
-----	----------	----------	----------	---

159	0.049302	0.001878	0.010146	C
-----	----------	----------	----------	---

160 rows × 16 columns



After we figure out the best parameters, we create a new DecisionTree Classifier and pass in the GridSearchCV results in order to produce the metrics and a confusion matrix for easier visualization.

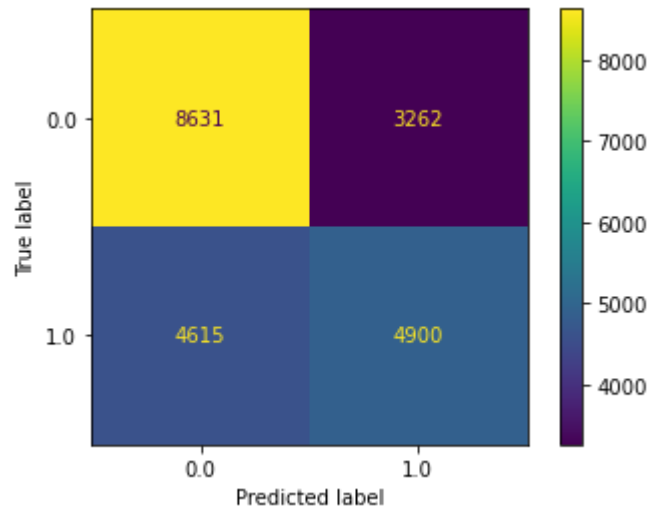
```
In [47]: dt2 = DecisionTreeClassifier(max_depth = 7, min_
dt2.fit(X_smote, y_smote)
y_dt2_pred = dt.predict(X_smote)
y_dt2_test_pred = dt.predict(X_test_smote)
dt2.score(X_test_smote, y_test_smote)
```

Out[47]: 0.6320534379671151

In [48]: `confusion_and_metrics(dt2, X_test_smote, y_test_s`

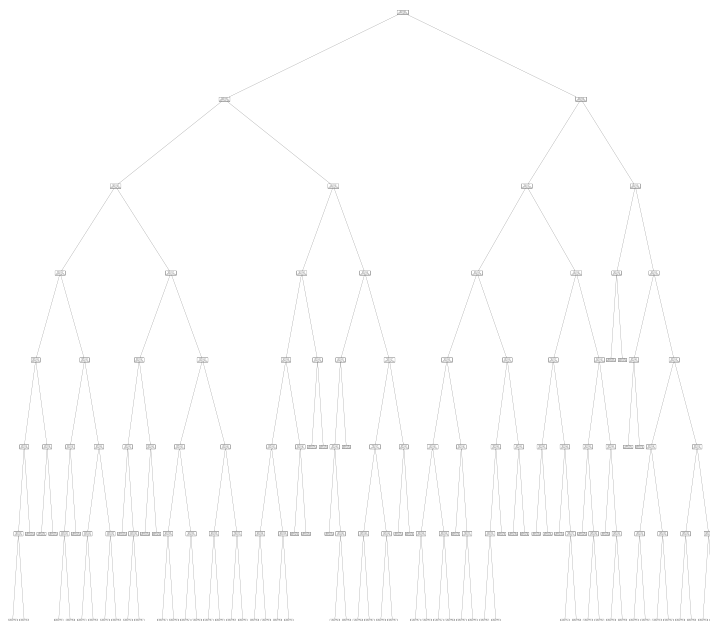
Accuracy Score: 0.632

Precision Score: 0.507



We also plotted the DecisionTree and exported it out to `out.pdf` for better visualization in a PDF Reader program.

In [49]: `f, ax = plt.subplots(figsize=(100, 100))  
plot_tree(dt2, ax=ax);  
# plt.savefig('out.pdf')`



From this DecisionTree, we recognize that our `X[21]` is one of the most important features for us to split our

one of the most important features for us to split our data. We exported the DecisionTree and took a closer look at the features to split on. Next we will run a LogisticRegression model.

### 3.3 Model 3 (Logistic Regression)

After running our DecisionTree Classifier, we implemented a LogisticRegression model to find our best predictions on H1N1 knowledge.

```
In [50]: # Importing the appropriate library
from sklearn.linear_model import LogisticRegression
model = LogisticRegression(random_state=42)
model.fit(X_smote, y_smote)
y_lr_pred = model.predict(X_smote)
y_lr_test_pred = model.predict(X_test_smote)
model.score(X_test_smote, y_test_smote)
```

```
C:\Users\Beter\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
0.6391068759342302
```

Out[50]:

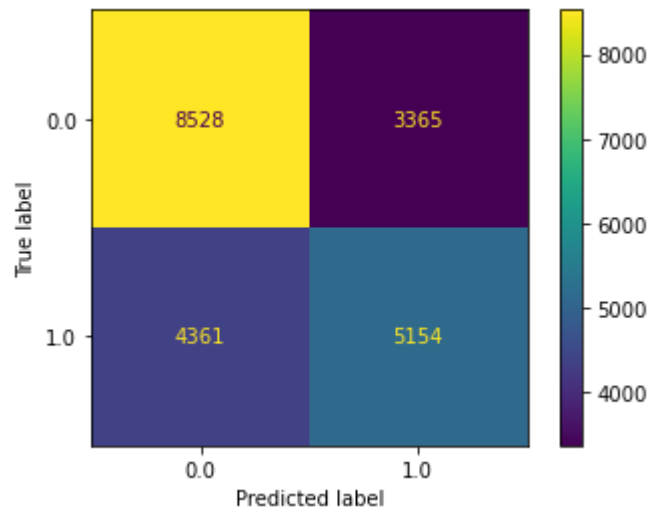
We plot out the confusion matrix and produce the metrics to see that our model is accurate in predicting whether a respondent is knowledgeable about H1N1 or not about **64%** of the time. This is a **14%** increase from our baseline model. The precision of this model also increased about **10%** meaning that our model correctly identifies knowledgeable respondents **60%** of the time.

The precision increase to 60% is important to us because in our models we would like to focus more on those who responded that they are knowledgeable about H1N1 Flu and Vaccine, but in reality they are not knowledgeable at all (*False Negative*).

```
In [51]: confusion_and_metrics(model, X_test_smote, y_test)
```

Accuracy Score: 0.639

Precision Score: 0.605



## 4. Results

### Interpreting LogisticRegression Results

We found that our LogisticRegression model produced the highest accuracy score of **64%** and the highest precision score of **60%**.

We want to also identify the coefficients in this array produced by our model. We want to identify the lowest coefficient and take the power of that coefficient in order to produce an odds value.

In the cell below, we are simply extracting the column names from the dataset, inputting them into a dictionary, and flipping the dictionary values. We then rename the columns in this coefficient dataframe to reflect the appropriate changes.

```
In [52]: # Extracting column names into a dictionary
model_column_names = {c: i for i, c in enumerate(model.columns)}

# Flipping the column keys and values
model_column_names_flipped = {model_column_names[c]: c for c in model_column_names}

# Turning the coefficients array into a dataframe
model_coef = pd.DataFrame(model.coef_)
```



```

model_coef.rename(model_column_names_flipped, axis=1)
# Checking to see if the rename was done correctly
model_coef.columns
# Those with an education of less than 12 years

```

```

Out[52]: Index(['h1n1_concern', 'behavioral_antiviral_meds', 'behavioral_avoidance',
        'behavioral_face_mask', 'behavioral_wash_hands',
        'behavioral_large_gatherings', 'behavioral_outside_home',
        'behavioral_touch_face', 'doctor_recc_h1n1', 'doctor_recc_seasonal',
        'chronic_med_condition', 'opinion_h1n1_vacc_effective',
        'opinion_h1n1_risk', 'opinion_h1n1_sick_from_vacc',
        'opinion_seas_vacc_effective', 'age_group_18 - 34 Years',
        'age_group_35 - 44 Years', 'age_group_45 - 54 Years',
        'age_group_55 - 64 Years', 'age_group_65+ Years', 'education_12 Years',
        'education_< 12 Years', 'education_College Graduate',
        'education_Some College', 'race_Black', 'race_Hispanic',
        'race_Other or Multiple', 'race_White', 'sex_Female', 'sex_Male',
        'income_poverty_<= $75,000, Above Poverty', 'income_poverty_> $75,000',
        'income_poverty_Below Poverty'],
        dtype='object')

```

```

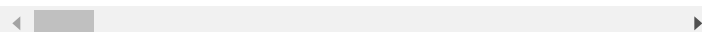
In [53]: # Printing out new dataframe to check if renaming worked
model_coef

```

```

Out[53]:
   h1n1_concern  behavioral_antiviral_meds  behavioral_avoidance
0      0.034431                0.083164                0.083164

```



We identify that from this array, the column of `education_< 12 Years` was our lowest coefficient. To make sense of this number we take the exponent of this value.

```

In [54]: # Calling function that defined above to interpret coefficients
print_odds(model_coef, 'education_< 12 Years')

```

```

education_< 12 Years: -0.7209944344950889
Odds: 0.4862684533209686

```

From this we see that those with an education\_< 12 Years are about **0.49x** as likely to be knowledgeable about the H1N1 flu and vaccine.

```
In [55]: print_odds(model_coef, 'income_poverty_Below Poverty')
```

```
income_poverty_Below Poverty: -0.34013342536524627  
Odds: 0.7116753608826263
```

From this we can see that respondents who indicate income\_poverty\_Below Poverty are about **0.69x** as likely to be knowledgeable about the H1N1 Flu and Vaccine.

```
In [56]: print_odds(model_coef, 'sex_Male')
```

```
sex_Male: -0.12456103754841111  
Odds: 0.8828843706242121
```

From this we can see that respondents who indicate sex\_Male are about **0.85x** as likely to be knowledgeable about the H1N1 Flu and Vaccine.

## Visualizations

Below we will be graphing some of the important features that can lead to misinformation: **overall knowledge of H1N1** and **concern for H1N1** as they both can lead to misinformation.

```
In [57]: def change_width(ax, new_value) :
```