

 [main](#) ▾



[Phase\\_3\\_Project](#) / [Project 3 Final Notebook.ipynb](#)



[petercvuong](#) Adjusted some markdown

 History

 1 contributor

1.3 MB





Centers for Disease  
Control and Prevention

## Contents

- 1. Business Understanding
- 1. EDA and Data Cleaning
- 1. Modeling
  - 3.1 Dummy Model
  - 3.2 DecisionTree Classifier
  - 3.3 LogisticRegression Classifier
- 1. Results
- 1. Conclusion
- 1. Recommendations
- 1. Future Considerations

## 1. Business Understanding

### Tackling Flu and Vaccine Misinformation for the CDC

The mission statement of the CDC is **"to promote health and quality of life by preventing and controlling disease, injury, and disability."**

And so, an integral part of fulfilling this mission is to provide the population with the appropriate information that is needed in order to make the best informed decisions for their health.

Disease and vaccine misinformation are a major

hurdle for the CDC. Lack of information or knowledge of the wrong information can lead to misinformed decisions that will compromise not only a single individual's health and safety but the population's as well.

The main goal of our project is to predict what features are most influential in determining a respondent's knowledge of the H1N1 flu and vaccine. Once we identify what features are most influential to a respondent's knowledge, we can focus on these features as areas of that the CDC should focus on in order to minimize the dangers of misinformation.

## 2. EDA and Data Cleaning

### Data and Limitations

The data that we will be working with in this project comes from the [DrivenData](#) website. This data is a survey response from 2009.

Since this data is from a survey response, it does not give an accurate reflection of the situation and community at the time. The data also has a pretty obvious class imbalance and some clear biases.

### Loading in the Data and some Initial EDA

We are looking at the vaccine data from 2009 about H1N1 flu and vaccine awareness.

The data survey response data that we got from DrivenData had about 27,000 recorded responses.

```
In [1]: # Importing the appropriate libraries that will be used
import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
from sklearn.feature_selection import RFE
from sklearn.preprocessing import PolynomialFeatures
import statsmodels
```

```

from statsmodels.formula.api import ols
from sklearn.dummy import DummyRegressor

from scipy import stats
from sklearn.preprocessing import OneHotEncoder
from folium.plugins import FastMarkerCluster

from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, recall_score
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import roc_auc_score, plot_roc_curve

```

In [2]:

```

# Loading in the pre-split datasets that were given
vaccinetrainingdf = pd.read_csv("data/training_set_features.csv")
vaccinetestdf = pd.read_csv("data/test_set_features.csv")
vaccinelabelsdf = pd.read_csv("data/training_set_labels.csv")

```

In [3]:

```

# Initial checking to see what data types we are working with
vaccinetrainingdf.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 36 columns):
 #   Column                                     Non-Null Count  Dtype  
---  --
 0   respondent_id                             26707 non-null  int64  
 1   h1n1_concern                             26615 non-null  float64
 2   h1n1_knowledge                           26591 non-null  float64
 3   behavioral_antiviral_meds                 26636 non-null  float64
 4   behavioral_avoidance                      26499 non-null  float64
 5   behavioral_face_mask                     26688 non-null  float64
 6   behavioral_wash_hands                     26665 non-null  float64
 7   behavioral_large_gatherings               26620 non-null  float64
 8   behavioral_outside_home                   26625 non-null  float64
 9   behavioral_touch_face                     26579 non-null  float64
10   doctor_recc_h1n1                         24547 non-null  float64
11   doctor_recc_seasonal                     24547 non-null  float64
12   chronic_med_condition                    25736 non-null  float64
13   child_under_6_months                     25887 non-null  float64

```

```

14 health_worker          25903 non-null
float64
15 health_insurance      14433 non-null
float64
16 opinion_h1n1_vacc_effective 26316 non-null
float64
17 opinion_h1n1_risk       26319 non-null
float64
18 opinion_h1n1_sick_from_vacc 26312 non-null
float64
19 opinion_seas_vacc_effective 26245 non-null
float64
20 opinion_seas_risk        26193 non-null
float64
21 opinion_seas_sick_from_vacc 26170 non-null
float64
22 age_group              26707 non-null
object
23 education              25300 non-null
object
24 race                   26707 non-null
object
25 sex                    26707 non-null
object
26 income_poverty         22284 non-null
object
27 marital_status         25299 non-null
object
28 rent_or_own            24665 non-null
object
29 employment_status      25244 non-null
object
30 hhs_geo_region         26707 non-null
object
31 census_msa             26707 non-null
object
32 household_adults       26458 non-null
float64
33 household_children     26458 non-null
float64
34 employment_industry    13377 non-null
object
35 employment_occupation  13237 non-null
object
dtypes: float64(23), int64(1), object(12)
memory usage: 7.3+ MB

```

As we can see there are some missing values, and some of the values are objects. Since some of the values in these columns are objects, we know that we have to one hot encode the values in order to implement them mathematically into our models.

```
In [4]: vaccinetrainingdf.head()
```

```
Out[4]: respondent_id  h1n1_concern  h1n1_knowledge  behavior
```

<b>0</b>	0	1.0	0.0
<b>1</b>	1	3.0	2.0
<b>2</b>	2	1.0	1.0
<b>3</b>	3	1.0	1.0
<b>4</b>	4	2.0	1.0

5 rows × 36 columns

In [5]:

```
# We see here that the data is already split almost
vaccinetestdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26708 entries, 0 to 26707
Data columns (total 36 columns):
 #   Column                                Non-Null Count
Dtype  -----
-----
 0   respondent_id                        26708 non-null
int64
 1   h1n1_concern                        26623 non-null
float64
 2   h1n1_knowledge                      26586 non-null
float64
 3   behavioral_antiviral_meds           26629 non-null
float64
 4   behavioral_avoidance                26495 non-null
float64
 5   behavioral_face_mask                26689 non-null
float64
 6   behavioral_wash_hands                26668 non-null
float64
 7   behavioral_large_gatherings         26636 non-null
float64
 8   behavioral_outside_home              26626 non-null
float64
 9   behavioral_touch_face                26580 non-null
float64
10   doctor_recc_h1n1                    24548 non-null
float64
11   doctor_recc_seasonal                 24548 non-null
float64
12   chronic_med_condition                25776 non-null
float64
13   child_under_6_months                 25895 non-null
float64
14   health worker                        25919 non-null
```

```

float64      -
   15  health_insurance      14480 non-null
float64
   16  opinion_h1n1_vacc_effective  26310 non-null
float64
   17  opinion_h1n1_risk      26328 non-null
float64
   18  opinion_h1n1_sick_from_vacc  26333 non-null
float64
   19  opinion_seas_vacc_effective  26256 non-null
float64
   20  opinion_seas_risk      26209 non-null
float64
   21  opinion_seas_sick_from_vacc  26187 non-null
float64
   22  age_group      26708 non-null
object
   23  education      25301 non-null
object
   24  race      26708 non-null
object
   25  sex      26708 non-null
object
   26  income_poverty      22211 non-null
object
   27  marital_status      25266 non-null
object
   28  rent_or_own      24672 non-null
object
   29  employment_status      25237 non-null
object
   30  hhs_geo_region      26708 non-null
object
   31  census_msa      26708 non-null
object
   32  household_adults      26483 non-null
float64
   33  household_children      26483 non-null
float64
   34  employment_industry      13433 non-null
object
   35  employment_occupation      13282 non-null
object
dtypes: float64(23), int64(1), object(12)
memory usage: 7.3+ MB

```

In [6]:

```

# Utilizing this block of code just to display a
pd.set_option('max_columns', None)
vaccinetestdf.head()

```

Out[6]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavior
0	26707	2.0	2.0	
1	26708	1.0	1.0	
2	26709	2.0	2.0	

2	26709	2.0	2.0
3	26710	1.0	1.0
4	26711	3.0	1.0

In [7]:

```
vaccinelabelsdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   respondent_id    26707 non-null  int64
1   h1n1_vaccine     26707 non-null  int64
2   seasonal_vaccine 26707 non-null  int64
dtypes: int64(3)
memory usage: 626.1 KB
```

Next, we're going to be doing some data cleaning to get rid of any extraneous columns that we determined were not relevant to our business problem of tackling misinformation in flu and vaccine awareness.

In [8]:

```
# Dropping data we deemed unnecessary and irrelevant
columns_to_drop = ['respondent_id', 'h1n1_knowledge',
                   'employment_status', 'rent_or_ownership',
                   'health_worker']
X_train = vaccinetrainingsdf.copy().drop(columns_to_drop)
X_test = vaccinetestdf.copy().drop(columns_to_drop)

# Setting the y_train and y_test to just be the h1n1_knowledge
y_train = vaccinetrainingsdf['h1n1_knowledge']
y_test = vaccinetestdf['h1n1_knowledge']
```

In [9]:

```
# Checking our data again to see that we dropped
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 20 columns):
#   Column          Non-Null Count  Dtype
---  -
0   h1n1_concern     26615 non-null  float64
1   behavioral_antiviral_meds 26636 non-null  float64
2   behavioral_avoidance 26499 non-null
```



```

float64
 3  behavioral_face_mask      26688 non-null
float64
 4  behavioral_wash_hands     26665 non-null
float64
 5  behavioral_large_gatherings 26620 non-null
float64
 6  behavioral_outside_home    26625 non-null
float64
 7  behavioral_touch_face     26579 non-null
float64
 8  doctor_recc_h1n1          24547 non-null
float64
 9  doctor_recc_seasonal      24547 non-null
float64
10  chronic_med_condition     25736 non-null
float64
11  opinion_h1n1_vacc_effective 26316 non-null
float64
12  opinion_h1n1_risk          26319 non-null
float64
13  opinion_h1n1_sick_from_vacc 26312 non-null
float64
14  opinion_seas_vacc_effective 26245 non-null
float64
15  age_group                 26707 non-null
object
16  education                 25300 non-null
object
17  race                      26707 non-null
object
18  sex                       26707 non-null
object
19  income_poverty            22284 non-null
object
dtypes: float64(15), object(5)
memory usage: 4.1+ MB

```

In [10]:

```
X_test.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26708 entries, 0 to 26707
Data columns (total 20 columns):
#   Column                                Non-Null Count
Dtype
---  ---
-----
0   h1n1_concern                          26623 non-null
float64
1   behavioral_antiviral_meds            26629 non-null
float64
2   behavioral_avoidance                 26495 non-null
float64
3   behavioral_face_mask                 26689 non-null
float64
4   behavioral_wash_hands                26668 non-null
float64
5   behavioral_large_gatherings          26636 non-null
float64
6   behavioral outside home              26626 non-null

```

```

float64
 7  behavioral_touch_face      26580 non-null
float64
 8  doctor_recc_h1n1          24548 non-null
float64
 9  doctor_recc_seasonal      24548 non-null
float64
10  chronic_med_condition      25776 non-null
float64
11  opinion_h1n1_vacc_effective 26310 non-null
float64
12  opinion_h1n1_risk           26328 non-null
float64
13  opinion_h1n1_sick_from_vacc 26333 non-null
float64
14  opinion_seas_vacc_effective 26256 non-null
float64
15  age_group                  26708 non-null
object
16  education                  25301 non-null
object
17  race                       26708 non-null
object
18  sex                        26708 non-null
object
19  income_poverty             22211 non-null
object
dtypes: float64(15), object(5)
memory usage: 4.1+ MB

```

## Feature Engineering

Created some frequently used functions that we will be utilizing throughout our project

```

In [11]: # Defined a OneHotEncoder function for ease of use
def OHE(X_train, categories):
    onehot = OneHotEncoder(sparse=False, handle_unknown='ignore')
    x_train_cat = pd.DataFrame(onehot.fit_transform(X_train[categories]).toarray())
    x_train_cat.columns = onehot.get_feature_names_out(categories)

    # Reset indices to avoid merging conflicts
    x_train_cat.reset_index(drop=True, inplace=True)
    X_train.reset_index(drop=True, inplace=True)

    # Joined the OHE dataframe to the dataframe
    x_train_df = X_train.drop(columns=categories, axis=1).join(x_train_cat)
    return x_train_df

# Defined a function that takes in parameters to calculate metrics
def confusion_and_metrics(model, X_test, y_test):
    # Accuracy Score
    print(f"Accuracy Score: {model.score(X_test, y_test)}")

    # Precision Score
    print(f"Precision Score: {precision_score(y_test, model.predict(X_test))}")

```

```

# Plot confusion matrix for visualization
plot_confusion_matrix(model, X_test, y_test)

# Defined a function to take in column name and dataframe
def print_odds(dataframe, column_name):
    # Prints out the name of the column and it's odds
    print(f"{column_name}: {dataframe[column_name].value_counts()}")

    # Prints out the odds value of the column
    print(f"Odds: {np.exp(dataframe[column_name].value_counts().values)}")

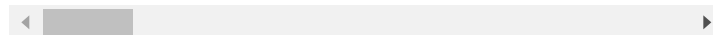
```

In [12]: X\_train

Out[12]:

	h1n1_concern	behavioral_antiviral_meds	behavioral_concern
0	1.0	0.0	0.0
1	3.0	0.0	0.0
2	1.0	0.0	0.0
3	1.0	0.0	0.0
4	2.0	0.0	0.0
...	...	...	...
26702	2.0	0.0	0.0
26703	1.0	0.0	0.0
26704	2.0	0.0	0.0
26705	1.0	0.0	0.0
26706	0.0	0.0	0.0

26707 rows × 20 columns



## SimpleImputer to Account for NaN Values

Prior to running some classification models on our data, we looked at it again and noticed that there were still a couple missing values.

In order to rectify this, we created a simple imputer to replace the NaN values with the most frequent

value(otherwise known as the mode) in its respective column.

We chose to use the mode to replace these NaN values because using the mode will keep the distribution of the data consistent.

```
In [13]: # Created a SimpleImputer to replace the NaN values  
from sklearn.impute import SimpleImputer  
imputer = SimpleImputer(strategy = 'most_frequent')  
imputed_X_train = imputer.fit_transform(X_train)  
imputed_X_train_df = pd.DataFrame(imputed_X_train)
```

```
In [14]: imputed_X_test = imputer.transform(X_test)  
imputed_X_test_df = pd.DataFrame(imputed_X_test)  
imputed_X_test_df
```

```
Out[14]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13
0	2	0	1	0	1	1	0	1	0	0	0	5	1	1
1	1	0	0	0	0	0	0	0	0	0	0	4	1	1
2	2	0	0	1	1	1	1	1	0	0	0	5	4	2
3	1	0	0	0	0	0	0	0	1	1	1	4	2	2
4	3	1	1	0	1	1	1	1	0	0	0	5	2	4
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
26703	1	0	1	0	1	0	0	1	1	1	0	4	2	2
26704	3	0	1	0	1	1	1	1	0	0	0	4	1	1
26705	0	0	0	0	0	0	0	0	0	0	0	4	3	1
26706	3	0	1	0	1	0	1	0	0	0	0	2	3	4

26707 2 0 0 0 1 0 0 1 1 0 0 5 1 2

26708 rows × 20 columns

In [15]: `# After doing the imputation and renaming, check  
imputed_X_test_df.info()`

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 26708 entries, 0 to 26707  
Data columns (total 20 columns):  
#   Column  Non-Null Count  Dtype  
---  -  
0    0      26708 non-null    object  
1    1      26708 non-null    object  
2    2      26708 non-null    object  
3    3      26708 non-null    object  
4    4      26708 non-null    object  
5    5      26708 non-null    object  
6    6      26708 non-null    object  
7    7      26708 non-null    object  
8    8      26708 non-null    object  
9    9      26708 non-null    object  
10   10     26708 non-null    object  
11   11     26708 non-null    object  
12   12     26708 non-null    object  
13   13     26708 non-null    object  
14   14     26708 non-null    object  
15   15     26708 non-null    object  
16   16     26708 non-null    object  
17   17     26708 non-null    object  
18   18     26708 non-null    object  
19   19     26708 non-null    object  
dtypes: object(20)  
memory usage: 4.1+ MB
```

After imputing the data, we recognize that the column names have disappeared but the indices are still there. To resolve this, we create a dictionary with the original column names and call the rename function on this new data frame's columns.

In [16]: `# Extracting column names into a dictionary  
dictionary_of_names = {columns: index for index,  
  
# Flipping the column keys and values  
dictionary_of_names_flipped = {dictionary_of_name  
  
# Checking to see if the column names were extra  
dictionary_of_names_flipped`

```
Out[16]: {0: 'h1n1_concern',
          1: 'behavioral_antiviral_meds',
          2: 'behavioral_avoidance',
          3: 'behavioral_face_mask',
          4: 'behavioral_wash_hands',
          5: 'behavioral_large_gatherings',
          6: 'behavioral_outside_home',
          7: 'behavioral_touch_face',
          8: 'doctor_recc_h1n1',
          9: 'doctor_recc_seasonal',
          10: 'chronic_med_condition',
          11: 'opinion_h1n1_vacc_effective',
          12: 'opinion_h1n1_risk',
          13: 'opinion_h1n1_sick_from_vacc',
          14: 'opinion_seas_vacc_effective',
          15: 'age_group',
          16: 'education',
          17: 'race',
          18: 'sex',
          19: 'income_poverty'}
```

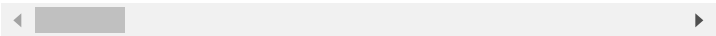
```
In [17]: # Created new variables for the training and test
         imputed_X_train_df_plus_column_names = imputed_X_
         imputed_X_test_df_plus_column_names = imputed_X_
```

```
In [18]: # Calling the new dataframe variables to check if
         imputed_X_test_df_plus_column_names
```

```
Out[18]:
```

	h1n1_concern	behavioral_antiviral_meds	behaviora
0	2		0
1	1		0
2	2		0
3	1		0
4	3		1
...	...		...
26703	1		0
26704	3		0
26705	0		0
26706	3		0
26707	2		0

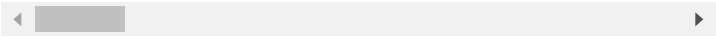
26708 rows × 20 columns



```
In [19]: imputed_X_train_df_plus_column_names
```

	h1n1_concern	behavioral_antiviral_meds	behaviora
0	1	0	
1	3	0	
2	1	0	
3	1	0	
4	2	0	
...	...	...	
26702	2	0	
26703	1	0	
26704	2	0	
26705	1	0	
26706	0	0	

26707 rows × 20 columns

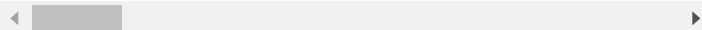


```
In [20]: X_train
```

	h1n1_concern	behavioral_antiviral_meds	behaviora
0	1.0	0.0	
1	3.0	0.0	
2	1.0	0.0	
3	1.0	0.0	
4	2.0	0.0	

...	...	...
<b>26702</b>	2.0	0.0
<b>26703</b>	1.0	0.0
<b>26704</b>	2.0	0.0
<b>26705</b>	1.0	0.0
<b>26706</b>	0.0	0.0

26707 rows × 20 columns



After doing some initial data cleaning and making sure that our data was uniform, we next want to address the problem of having the `object` type in our columns. In order to address this, we apply a `OneHotEncoder` to these object columns of `age_group`, `education`, `race`, `sex`, and `income_poverty`.

We check the values in each of these object columns to see how many variables will be OneHotEncoded.

```
In [21]: X_train['age_group'].value_counts()
```

```
Out[21]: 65+ Years      6843
55 - 64 Years    5563
45 - 54 Years    5238
18 - 34 Years    5215
35 - 44 Years    3848
Name: age_group, dtype: int64
```

```
In [22]: X_train['education'].value_counts()
```

```
Out[22]: College Graduate    10097
Some College      7043
12 Years          5797
< 12 Years        2363
Name: education, dtype: int64
```

```
In [23]: X_train['race'].value_counts()
```

```
Out[23]: White      21222
Black      2118
Hispanic   1755
Other or Multiple  1612
Name: race, dtype: int64
```



```
In [24]: X_train['sex'].value_counts()
```

```
Out[24]: Female    15858  
Male      10849  
Name: sex, dtype: int64
```

```
In [25]: X_train['income_poverty'].value_counts()
```

```
Out[25]: <= $75,000, Above Poverty    12777  
> $75,000                    6810  
Below Poverty                 2697  
Name: income_poverty, dtype: int64
```

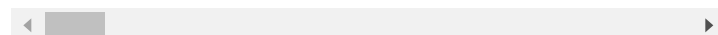
After counting the values, we see that **18** columns will be added

```
In [26]: # Called the OHE function we made and assigned new columns  
ohe_training_df = OHE(imputed_X_train_df_plus_columns, ohe_dictionary)  
ohe_test_df = OHE(imputed_X_test_df_plus_columns, ohe_dictionary)  
ohe_training_df
```

```
Out[26]:
```

	h1n1_concern	behavioral_antiviral_meds	behavioral_risk_factors
0	1	0	0
1	3	0	0
2	1	0	0
3	1	0	0
4	2	0	0
...	...	...	...
26702	2	0	0
26703	1	0	0
26704	2	0	0
26705	1	0	0
26706	0	0	0

26707 rows × 33 columns



We finished OneHotEncoding the object values and now we have to bin the target values.

Based on the [data dictionary](#), our target values reside in the `h1n1_knowledge` column where the responses are recorded as such:

- 0 = No knowledge
- 1 = A Little Knowledge
- 2 = A Lot of Knowledge

For our project, we are going to bin the 0s and 1s together because those who respond as having little to no knowledge of the flu and vaccine are most prone to misinformation.

We will then be turning all the 2 responses into 1s so that we have a simple binary categorization where:

- 0 = Little/No Knowledge
- 1 = A Lot of Knowledge

```
In [27]: # Instead of calling SimpleImputer and removing
# imputation which replaced all the NaN values with
# which in this case would be 1.0 (Little knowledge)
y_train.replace(np.nan, 1.0, inplace = True)
y_test.replace(np.nan, 1.0, inplace = True)
```

```
In [28]: # Checking to see if we replaced the NaN values
y_test.isna().value_counts()
```

```
Out[28]: False      26708
Name: h1n1_knowledge, dtype: int64
```

```
In [29]: y_train.isna().value_counts()
```

```
Out[29]: False      26707
Name: h1n1_knowledge, dtype: int64
```

```
In [30]: # Binning all the 1.0s with the 0.0s
y_train.replace(1.0, 0.0, inplace = True)
y_test.replace(1.0, 0.0, inplace = True)
# Replacing all the 2.0s with 1.0s
y_train.replace(2.0, 1.0, inplace = True)
y_test.replace(2.0, 1.0, inplace = True)
```

```
In [31]: # Checking to see if we replaced our values correctly
y_train.value_counts()
```

```
Out[31]: 0.0      17220
1.0       9487
Name: h1n1_knowledge, dtype: int64
```

```
In [32]: y_test.value_counts()
```

```
Out[32]: 0.0    17193
          1.0     9515
          Name: h1n1_knowledge, dtype: int64
```

```
In [33]: ohe_training_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 33 columns):
 #   Column                                No
n-Null Count  Dtype
---  -
-----
 0   h1n1_concern                          26
707 non-null object
 1   behavioral_antiviral_meds             26
707 non-null object
 2   behavioral_avoidance                  26
707 non-null object
 3   behavioral_face_mask                  26
707 non-null object
 4   behavioral_wash_hands                  26
707 non-null object
 5   behavioral_large_gatherings            26
707 non-null object
 6   behavioral_outside_home                26
707 non-null object
 7   behavioral_touch_face                  26
707 non-null object
 8   doctor_recc_h1n1                      26
707 non-null object
 9   doctor_recc_seasonal                  26
707 non-null object
10   chronic_med_condition                  26
707 non-null object
11   opinion_h1n1_vacc_effective             26
707 non-null object
12   opinion_h1n1_risk                       26
707 non-null object
13   opinion_h1n1_sick_from_vacc             26
707 non-null object
14   opinion_seas_vacc_effective             26
707 non-null object
15   age_group_18 - 34 Years                26
707 non-null float64
16   age_group_35 - 44 Years                26
707 non-null float64
17   age_group_45 - 54 Years                26
707 non-null float64
18   age_group_55 - 64 Years                26
707 non-null float64
19   age_group_65+ Years                    26
707 non-null float64
20   education_12 Years                     26
707 non-null float64
21   education_< 12 Years                   26
707 non-null float64
22   education_College Graduate             26
707 non-null float64
23   education_Some College                  26
```

```

23 education_some_college
707 non-null float64
24 race_Black
707 non-null float64
25 race_Hispanic
707 non-null float64
26 race_Other or Multiple
707 non-null float64
27 race_White
707 non-null float64
28 sex_Female
707 non-null float64
29 sex_Male
707 non-null float64
30 income_poverty_<= $75,000, Above Poverty
707 non-null float64
31 income_poverty_> $75,000
707 non-null float64
32 income_poverty_Below Poverty
707 non-null float64
dtypes: float64(18), object(15)
memory usage: 6.7+ MB

```

## SMOTE for Class Imbalance

After we bin our target and features together, we recognize that our target class is severely imbalanced. To address this class imbalance, we implement SMOTE to undersample our 0 class.

```

In [34]: # Since our data is severely imbalanced, we utilize
# Since we SMOTE our training dataset, we must SM

from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
from collections import Counter

over = SMOTE(sampling_strategy=0.7)
under = RandomUnderSampler(sampling_strategy=0.8)

X_smote, y_smote = under.fit_resample(ohe_traini
X_test_smote, y_test_smote = under.fit_resample(

counter = Counter(y_train)
test_counter = Counter(y_test_smote)
print(counter)
print(test_counter)

```

```

Counter({0.0: 17220, 1.0: 9487})
Counter({0.0: 11893, 1.0: 9515})

```

## Checking for Preliminary Feature Importance

For the final part of our EDA and Data cleaning, we want to check and see what features are seemingly

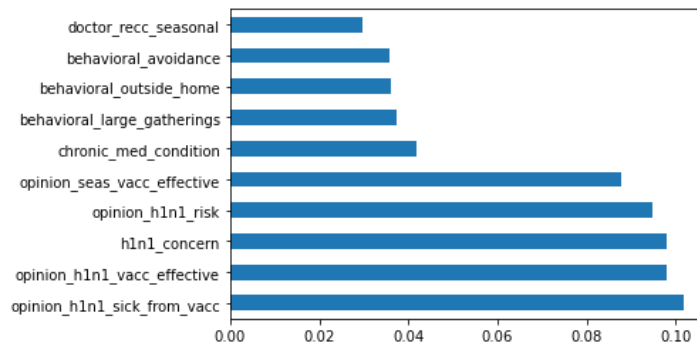
most important to our respondents.

```
In [35]: # Feature columns
X = ohe_training_df.iloc[:,0:33]
# Target column - H1N1 Knowledge
y = vaccinetrainingdf.iloc[:,2]

from sklearn.ensemble import ExtraTreesClassifier
import matplotlib.pyplot as plt

# Instantiate model
modelfeatures = ExtraTreesClassifier()
modelfeatures.fit(X,y)
print(modelfeatures.feature_importances_) # use l
# Plot graph of feature importances for better v
feat_importances = pd.Series(modelfeatures.feature
feat_importances.nlargest(10).plot(kind='barh')
plt.show()
```

```
[0.09806714 0.01422531 0.03588687 0.01741775 0.02
203115 0.03724437
 0.03613427 0.02670774 0.02055141 0.02978196 0.04
196037 0.09807479
 0.0947744 0.10175311 0.08761582 0.01497825 0.01
578261 0.01660707
 0.01707808 0.01495026 0.01276786 0.01256776 0.02
436088 0.00960014
 0.00782002 0.00751059 0.00828499 0.01456991 0.01
27635 0.01242991
 0.01115611 0.01591016 0.00863544]
```



We notice that the top 5 features that are most important to survey respondents are:

- opinion\_h1n1\_sick\_from\_vacc
- opinion\_h1n1\_vacc\_effective
- h1n1\_concern
- opinion\_h1n1\_risk
- opinion\_seas\_vacc\_effective

### 3. Modeling

Now that our data is cleaned, we can go into the

modeling.

### 3.1 Model 1 (Dummy Classifier)

First we want to create a DummyClassifier model that will serve as the baseline for our model performance comparison. A DummyClassifier model in this case would mean that based on the given data, the dummy model would correctly identify our predictions **50%** of the time.

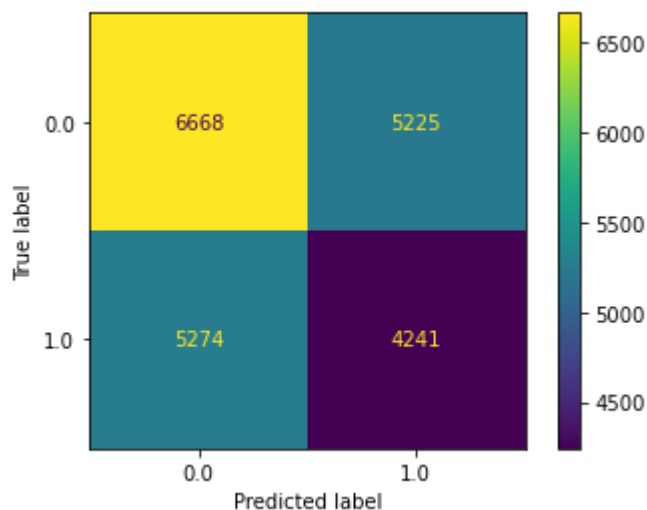
```
In [36]: # Created Dummy Classifier model to look at simple
from sklearn.dummy import DummyClassifier
dummy = DummyClassifier()
dummy.fit(X_smote, y_smote)
y_pred = dummy.predict(X_smote)
y_test_pred = dummy.predict(X_test_smote)
y_pred_df = pd.DataFrame(y_pred)
dummy.score(X_test_smote, y_test_smote)
```

```
C:\Users\Beter\anaconda3\envs\learn-env\lib\site-
packages\sklearn\dummy.py:131: FutureWarning: The
default value of strategy will change from strati
fied to prior in 0.24.
  warnings.warn("The default value of strategy wi
ll change from "
0.5039237668161435
```

Out[36]:

```
In [37]: # Called function and printed out confusion and
confusion_and_metrics(dummy, X_test_smote, y_test_smote)
```

Accuracy Score: 0.509  
Precision Score: 0.450



We print out the `dummy.score` to see that the accuracy score is about 50%, just as we expected.

We call our `confusion_and_metrics` function that we defined above in order to produce the evaluation metrics of Accuracy and Precision and a confusion matrix for easier visualization.

## 3.2 Model 2 (Decision Tree Classifier)

Next we will create an inferential DecisionTree Classifier in order to identify our **most important features**. After we identify our most important features, we can then run a classifying LogisticRegression model to measure our predictions on the dataset.

```
In [38]: # Displayed cross validation score for the dummy
from sklearn.model_selection import train_test_split

dummy_cross_val = cross_val_score(dummy, X_smote, y_smote, cv=5)

print(dummy_cross_val)
```

```
C:\Users\Beter\anaconda3\envs\learn-env\lib\site-packages\sklearn\dummy.py:131: FutureWarning: The default value of strategy will change from stratified to prior in 0.24.
  warnings.warn("The default value of strategy will change from "
Out[38]: array([0.49988288, 0.509487 , 0.50714453, 0.51487468, 0.51089248])
```

```
In [39]: from sklearn.tree import DecisionTreeClassifier,
from sklearn.model_selection import GridSearchCV

dt = DecisionTreeClassifier (random_state = 10)
dt.fit(X_smote, y_smote)
y_dt_pred = dt.predict(X_smote)
y_dt_test_pred = dt.predict(X_test_smote)
dt.score(X_test_smote, y_test_smote)
```

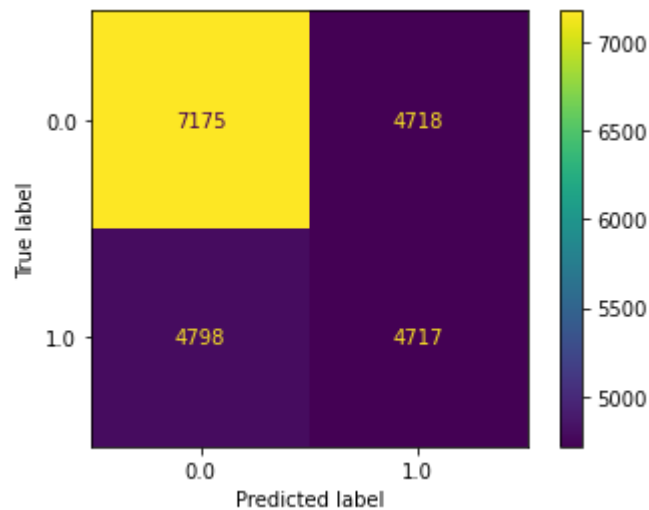
```
Out[39]: 0.5554932735426009
```

Similarly like how we checked the DummyClassifier, we also check the plot the confusion matrix and check the metrics of our baseline decision tree.

```
In [40]: # Called function and printed out confusion matrix
confusion_and_metrics(dt, X_test_smote, y_test_smote)
```

Accuracy Score: 0.555

Precision Score: 0.500



After running our initial DecisionTree Classifier, we got an accuracy score of about **56%**. As you can see, this accuracy score is only about 6% better than the baseline.

## Implementing GridSearchCV to Find Optimal Hyperparameters

We decided to use a GridSearchCV in order find the best hyperparameters to pass into our DecisionTree Classifier so that we can find the most important features to focus on.

```
In [41]: # Created grid parameter to perform a GridSearch
grid = {
    'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'min_samples_split': [100, 500, 1000, 5000],
    'min_samples_leaf': [100, 500, 1000, 5000]
}
# initializing our grid search with the grid parameter
gs = GridSearchCV(estimator = dt, param_grid = grid)
gs.fit(ohe_training_df, y_train)
```

```
Out[41]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(
    random_state=10),
    param_grid={'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'min_samples_leaf': [100, 500, 1000, 5000],
    'min_samples_split': [100, 500, 1000, 5000]})
```

After we run our GridSearchCV, we print out the best parameters, best score, and the



`best_params_`, `best_score_`, and the `best_estimator_` to get the optimal parameters and metrics based on the grid search results.

In [42]: `gs.best_params_`

Out[42]: `{'max_depth': 7, 'min_samples_leaf': 100, 'min_samples_split': 100}`

In [43]: `gs.best_score_`

Out[43]: `0.6770884108866996`

In [44]: `gs.best_estimator_.score(ohe_test_df, y_test)`

Out[44]: `0.6696495432080276`

In [45]: `gs.n_features_in_`

Out[45]: `33`

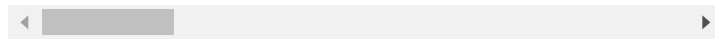
In [46]: `# Converted the results from the GridSearch to a pd.DataFrame(gs.cv_results_)`

Out[46]:

	mean_fit_time	std_fit_time	mean_score_time	std_sco
0	0.048013	0.001324	0.009579	C
1	0.048081	0.002682	0.009973	C
2	0.049832	0.003685	0.010999	C
3	0.044780	0.000810	0.009999	C
4	0.044751	0.000749	0.009766	C
...	...	...	...	
155	0.051214	0.000875	0.008188	C
156	0.044173	0.000333	0.008193	C

<b>157</b>	0.044093	0.000408	0.008498	C
<b>158</b>	0.044443	0.001037	0.008314	C
<b>159</b>	0.046540	0.002286	0.008771	C

160 rows × 16 columns



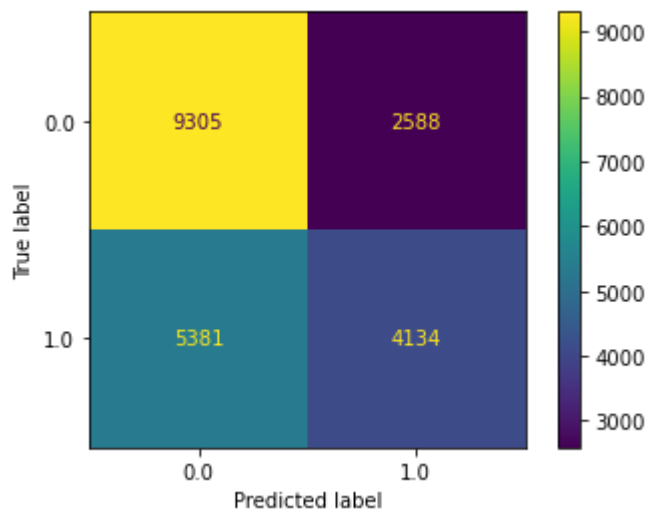
After we figure out the best parameters, we create a new DecisionTree Classifier and pass in the GridSearchCV results in order to produce the metrics and a confusion matrix for easier visualization.

```
In [47]: dt2 = DecisionTreeClassifier(max_depth = 7, min_
dt2.fit(X_smote, y_smote)
y_dt2_pred = dt.predict(X_smote)
y_dt2_test_pred = dt.predict(X_test_smote)
dt2.score(X_test_smote, y_test_smote)
```

Out[47]: 0.6277559790732437

```
In [48]: confusion_and_metrics(dt2, X_test_smote, y_test_
```

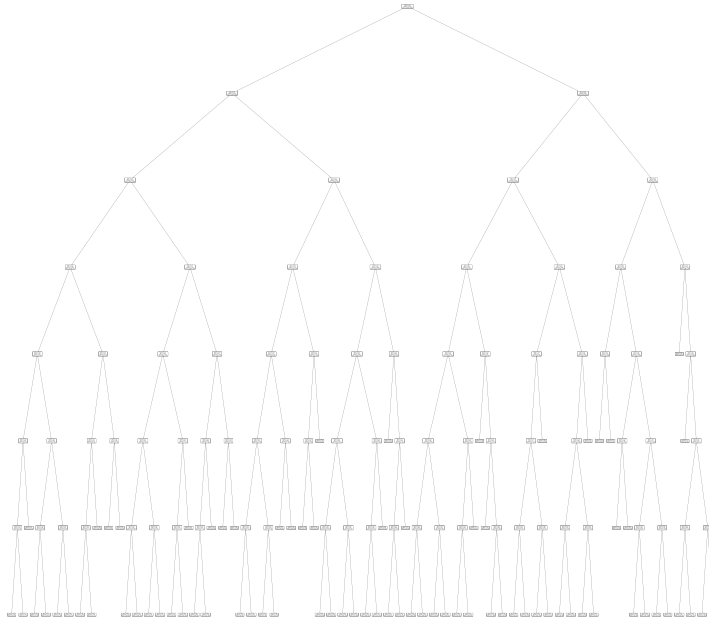
Accuracy Score: 0.628  
Precision Score: 0.500



We also plotted the DecisionTree and exported it out to out.pdf for better visualization in a PDF Reader program.

```
In [49]: f, ax = plt.subplots(figsize=(100, 100))
```

```
plot_tree(dt2, ax=ax);
# plt.savefig('out.pdf')
```



From this DecisionTree, we recognize that our  $X[21]$  is one of the most important features for us to split our data. We exported the DecisionTree and took a closer look at the features to split on. Next we will run a LogisticRegression model.

### 3.3 Model 3 (Logistic Regression)

After running our DecisionTree Classifier, we implemented a LogisticRegression model to find our best predictions on H1N1 knowledge.

```
In [50]: # Importing the appropriate library
from sklearn.linear_model import LogisticRegression
model = LogisticRegression(random_state=42)
model.fit(X_smote, y_smote)
y_lr_pred = model.predict(X_smote)
y_lr_test_pred = model.predict(X_test_smote)
model.score(X_test_smote, y_test_smote)
```

```
C:\Users\Beter\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

```

processing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
Out[50]: 0.6392470104633782

```

We plot out the confusion matrix and produce the metrics to see that our model is accurate in predicting whether a respondent is knowledgeable about H1N1 or not about **64%** of the time. This is a **14%** increase from our baseline model. The precision of this model also increased about **10%** meaning that our model correctly identifies knowledgeable respondents **60%** of the time.

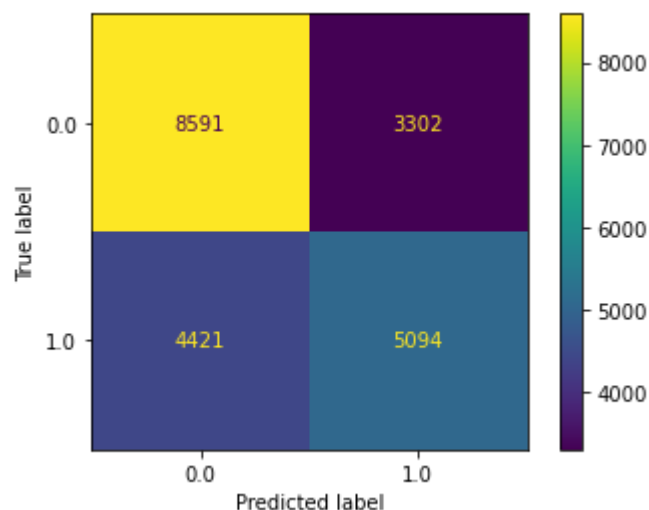
The precision increase to 60% is important to us because in our models we would like to focus more on those who responded that they are knowledgeable about H1N1 Flu and Vaccine, but in reality they are not knowledgeable at all (*False Negative*).

```

In [51]: confusion_and_metrics(model, X_test_smote, y_test)

Accuracy Score: 0.639
Precision Score: 0.607

```



## 4. Results

### Interpreting LogisticRegression Results

We found that our LogisticRegression model produced the highest accuracy score of **64%** and the highest precision score of **60%**.

We want to also identify the coefficients in this array produced by our model. We want to identify the lowest coefficient and take the power of that coefficient in order to produce an odds value.

In the cell below, we are simply extracting the column names from the dataset, inputting them into a dictionary, and flipping the dictionary values. We then rename the columns in this coefficient dataframe to reflect the appropriate changes.